



KinShield

The Complete Playbook

The full picture — explained for everyone, then for engineers and investors. The problem, the idea, the people, every feature and how it really works, the technology, the money, the step-by-step plan, the obstacles and their fixes, and where this grows into a category-defining company.

The orchestration layer for fraud defence in Canada.

One Trust Score behind every channel · built for the household · twelve languages · honest by engineering discipline.

START HERE

How to read this document

This is written so a non-technical reader and a CTO can both follow it. Throughout, you'll see four kinds of coloured boxes:

IN PLAIN WORDS

The simple version — no jargon. If you only read these, you'll understand the whole company.

HOW IT WORKS

The technical version — the engineering and the real-world mechanism.

WHY IT MATTERS

The business reason — what it wins us and why it's hard to copy.

REAL EXAMPLE

A concrete story of the feature catching a real scam.

Status tags on features: **BUILT** already working in the app today **DEMO** faithful interactive demo; needs telephony/OS/partner to go live **PLANNED** on the roadmap.

Every dollar figure is a planning estimate in CAD at realistic 2026 Canadian rates — deliberately conservative. Nothing here claims a capability the operating systems or the law don't permit; that honesty is the strategy (Book VIII).

Table of contents

BOOK I · The big picture (for everyone)

- 1 The problem, in plain words
- 2 The idea — the bodyguard, not the receptionist
- 3 Who it's for — four people, four days
- 4 Why now — the window

BOOK II · The market

- 5 The economics of fraud in Canada
- 6 Bill C-15, explained simply
- 7 How big is this? (sizing with the math shown)
- 8 The competition, one by one
- 9 The six moats

BOOK III · The product (every feature, clearly)

- 10 The Trust Score — the heart of everything
- 11 What happens when you check something (step by step)
- 12 The 30 features, broken out one by one
- 13 The data flywheel — how it gets smarter

BOOK IV · The technology

- 14 System architecture & the split data-plane
- 15 The Trust Engine pipeline
- 16 The Line — how cloud call-screening works
- 17 The phone problem — what Apple/Android allow & our solution
- 18 Security, privacy & compliance

BOOK V · Building it — cost, time, people

- 19 What every feature costs & how long it takes
- 20 The obstacle in each feature & how it's resolved
- 21 The four-stage roadmap, step by step
- 22 Who to hire, and when

BOOK VI · The business

- 23 Pricing & packaging
- 24 Go-to-market — how we get users cheaply
- 25 The five-year financial model
- 26 Unit economics — does each customer make money?
- 27 Funding & free government money

BOOK VII · The future — how this grows huge

- 28 Growth levers (next 12 months)

- 29 Future features (beyond the 30)
- 30 New markets & new countries
- 31 New revenue: telecom, insurance, government, B2B
- 32 Becoming a platform — the trust layer
- 33 Who buys us — the exit landscape

BOOK VIII · Risks & realities

- 34 The risk register
- 35 What founders forget
- 36 Investor FAQ — the hard questions answered
- 37 Claims discipline & ethics

Appendices

- A–E Glossary · scam scripts the engine knows · sample SDK API · what's built & how to run · 90-day checklist

BOOK I

The big picture

Everything in this book is jargon-free. If you read nothing else, read this.

SECTION 1

The problem, in plain words

Scammers have gone professional and gone AI. They can now **clone a voice** from a few seconds of audio, **spoof a real phone number**, build a **fake bank website** in minutes, and run **scripts in any language**. The result: a grandmother gets a call that sounds exactly like her grandson in jail; a newcomer gets a CRA "arrest" threat in Punjabi; a small-business bookkeeper joins a video call with a deepfaked CEO and wires \$48,000.

IN PLAIN WORDS

The old advice — "just be careful" — no longer works, because the fakes are now too good for a normal person to spot. People don't need another warning; they need something that **checks for them** and **knows their family**.

Why existing tools don't solve it

- **They protect one channel.** A spam-call blocker does nothing about a scam text or a fake investment website.
- **They're built for one tech-savvy American.** They're English-first and single-user — but the victims are seniors, newcomers, and whole families across iPhone + Android.
- **They stop at "detected."** When the worst happens — money sent, images leaked — they hand you a ticket number.
- **They overclaim.** "96% accurate" sounds great until a novel scam gets through and trust collapses.

WHY IT MATTERS

The gap isn't "better detection on one channel." It's a **missing layer**: something that watches every channel at once, understands the household, speaks its language, decides what to do, and stays through the recovery. That layer is the company.

SECTION 2

The idea — the bodyguard, not the receptionist

Apple and Google now answer unknown calls and transcribe them. That's a **receptionist** — it tells you who's at the door. KinShield is the **bodyguard** — it *judges* who's at the door, decides whether to let them in, and protects the whole family while doing it.

IN PLAIN WORDS

KinShield gives every message, link, call, payment and meeting a single easy number — a **Trust Score from 1 to 100** — plus the top reasons and one tap to act. Green is safe, red is a threat. A 78-year-old can use it; it even reads the answer out loud, in their language.

One brain sits behind everything: paste a suspicious text, scan a QR on a parking meter, forward a "job offer," let an unknown caller be screened — same Trust Score, same plain explanation, same actions (block, report, alert the family, start recovery). And the household is the unit: an adult child can protect a parent's phone from across the country.

WHY IT MATTERS

"Detection" is becoming free and built-in. What stays scarce and valuable is **judgment** (one fused verdict), **family context**, **Canadian + language context**, and **recovery**. KinShield sells those four — none of which the giants can cheaply copy.

SECTION 3

Who it's for — four people, four days

Daljit, 74 — the protected senior · Surrey, BC · speaks Punjabi

A call says her SIN is suspended and police are coming. KinShield's screening line answers first, recognises the CRA-arrest script, and blocks it — her phone never even rings. If she'd answered, she taps one big button and the app, in Punjabi, says aloud: *"This is a scam. Hang up. Do not pay."* Her daughter gets an alert instantly.

Priya, 41 — the adult child / buyer · pays the bill, protects two parents

She sets up the Family Circle in 60 seconds, enrolls both parents, and sets a policy: ping me on transfers over \$500. She sees flagged events on their phones, sleeps better, and tells three friends. **She is the wedge customer.**

Minh, 29 — the newcomer · Toronto · job & immigration scams

An "IRCC" text threatens his study permit; a "recruiter" offers easy money. KinShield checks both in Vietnamese, explains exactly what's wrong, and verifies the "company" against Canadian registries. **Day-1 install** for his community.

Sarah, 38 — the small-business owner · 6 staff · approves wires

On a video call her "CEO" urgently asks to change a vendor's bank details. Wire Guard requires an out-of-band release — a 5-second selfie + spoken codeword + voice match on her own phone. The deepfake can't pass it. \$48,000 stays put.

WHY IT MATTERS

Four segments, one product, one Trust Score — and each segment recruits the next (the senior's kids, the newcomer's community, the owner's staff). That's how a wedge becomes a movement.

SECTION 4

Why now — the window

THE THREAT EXPLODED

AI voice cloning, number spoofing and instant fake sites made scams unspottable by humans. Losses are at record highs and rising ~300% since 2020.

THE LAW JUST CHANGED

Bill C-15 (March 2026) forces every Canadian bank to detect & prevent fraud and report it — creating an instant B2B buyer for exactly what KinShield builds.

IN PLAIN WORDS

The danger is at an all-time high and the government just made fraud-prevention mandatory for banks — at the same moment. A window like this lasts quarters, not years. First credible mover with Canadian + multilingual depth takes the category.

BOOK II

The market

The numbers, the law, the size of the prize, and exactly how we beat each giant.

The economics of fraud in Canada

FACT (VERIFIED)	FIGURE
Reported fraud losses to the CAFC, 2024	\$638–643M CAD
Increase since 2020	~300%
Share of fraud actually reported (CAFC estimate)	5–10% → true losses in the billions
Rank of fraud as a complaint at Canada's banking ombudsman	#1 category
Global losses tied to impersonation fraud (INTERPOL, Mar 2026)	>US\$400B

REAL EXAMPLE

The "Arup" case: a finance employee joined a video meeting with a fully deepfaked executive and colleagues, and wired roughly **US\$25M**. No detector "caught" it — the fakes were perfect. Only a *procedure* (out-of-band verification) could have stopped it. That's the kind of attack KinShield is built against.

Bill C-15, explained simply

IN PLAIN WORDS

A new Canadian law (passed March 2026) says: **banks must now actively prevent fraud**, ask your permission before turning on risky features, let you set limits, and **report fraud data to the government**. Banks suddenly need software that does all this — and most don't have it.

What the law forces banks to do — and how KinShield's SDK maps to it

BILL C-15 REQUIREMENT	KINSHIELD SDK COMPONENT
Maintain policies to detect & prevent consumer-targeted fraud	Trust Engine API (score any transaction/message/session)
Obtain express consent before enabling risky capabilities	Drop-in consent & cool-off screens
Let customers disable capabilities & adjust limits	White-label capability toggles & limit UX
Report fraud data to the FCAC	FCAC reporting pipeline (residency-clean)

WHY IT MATTERS

This converts a consumer app into a **B2B revenue engine** with a legally-created, deadline-driven buyer. Credit unions and digital banks first (shorter sales cycles), Big-5 later.

SECTION 7

How big is this? (the math shown)

LAYER	WHAT IT COUNTS	ESTIMATE (CAD)
TAM (total)	NA consumer + SMB fraud/identity protection	US\$10B+ , double-digit growth
SAM (serviceable)	Canadian households + SMBs + the bank/CU SDK market	~\$1.5–2.5B / yr
SOM (Yr 3 realistic)	~45K paid households + 8–15 SDK customers + insurance	\$8M ARR path → \$30–60M by Yr 5

IN PLAIN WORDS (THE MATH)

Canada has ~15M households. Capturing just **1%** as paying customers = 150K households. At a blended ~\$140/year that's ~**\$21M** from consumers alone — before a single bank deal, insurance dollar, or telecom partnership. The consumer business funds the far bigger B2B business.

The competition, one by one

COMPETITOR	BEST AT	WEAK SPOT	OUR PLAY
Apple iOS Call Screening	Default distribution, free, built-in, 9 languages	Screens for convenience not fraud; no family, no Canada, no recovery	Build on top via Live Caller ID; add the judgment + family + recovery layer
Google Pixel Scam Detection	Deepest on-device AI; cryptographic fake-call detection	Pixel-only, single-user, English, generic patterns, nothing after detection	Own the cross-platform household; out-localize; own the aftermath
Norton Genie	Brand trust; free AI scam checker at scale	Reactive (you must ask), English-first, no Canadian context, no family	Ask KinShield: proactive, 12-language, Canada-tuned, wired to actions
McAfee Scam Detector	Strong suite; "96% deepfake" marketing; huge enrollments	Accuracy theatre; English-first; no real call defence; US data	Match coverage; win on calibrated honesty + Canadian depth
Truecaller	World's largest caller-ID database	Built from non-consenting contact uploads — a PIPEDA/Law 25 liability here; reactive blocklists	Judge behaviour not number; consent-clean network that's legal in Canada
Hiya	Carrier/OEM network reach (Samsung, AT&T)	Invisible (carrier owns the user), US-centric, single-user, no aftermath	Take the Canadian carrier road first; own the family relationship they never have
Robokiller / YouMail	Original answer-bot honeypots	Calls only, single-user, US, aging, no languages	Beat a feature with a system — Line + Honeypot inside a full platform
Aura	Best all-in-one identity bundle (US); \$1M insurance	US only; scam detection is an add-on; slow reimbursement	Be Canada's Aura first; add live interception, Senior Mode, pre-evidenced insurance
Incogni / DeleteMe	Data-broker removal at scale	US-broker-centric; removal is the whole product; no detection	Statutory Law-25 deletion demands + closed loop into detection
Pindrop / Reality Defender	Best enterprise deepfake research	\$50K+ enterprise only; no consumer; detectors generalize poorly to new fakes	Out-position: multi-signal + procedure stop a perfect deepfake; sell the consumer layer they don't make

SECTION 9

The six moats

#	MOAT	IN PLAIN WORDS	WHY IT COMPOUNDS
1	The Canadian corpus	A library of real scams in 12 languages	Built with settlement agencies + honeypot; data nobody can buy, refreshed daily
2	The family graph	Whole households enrol together	Switching costs balloon at 4 members; every senior recruits more families
3	Regulatory fluency	Compliance turned into a product	Bill C-15 kit, FCAC pipeline, Law 25 filings — years of legal work baked into the SDK
4	Split-plane privacy	We hold ciphertext we can't read	Data-monetizing incumbents can't retrofit E2EE without breaking their model
5	The recovery operation	Real trained humans, in your language	An operations moat — unglamorous, impossible to fake, the source of the brand's stories
6	Calibrated honesty	We publish our accuracy and our misses	Costs nothing; competitors' marketing departments will never allow it

WHY IT MATTERS

A strategy that only beats competitors at their *weaknesses* loses the moment they fix one. These six beat them at their *strengths* — and four of them get stronger the bigger we get.

BOOK III

The product

The Trust Score, exactly what happens when you check something, and all 30 features one by one.

The Trust Score — the heart of everything

IN PLAIN WORDS

One number, 1 to 100. **100 = looks safe. 1 = almost certainly fraud.** Below it: the top three reasons in your language, and a recommended action. That's the whole product's vocabulary — "what did it score?" Green, amber, red. A child or a grandparent understands it instantly.


HOW IT WORKS

The score is produced by **five independent "signal families,"** then fused. Crucially, **no single one decides** — they combine, so a perfect deepfake (which beats one signal) still loses to the others. The score comes with an **uncertainty band** ($\pm$) so a thin-evidence check honestly says "I'm less sure."

SIGNAL FAMILY	WHAT IT READS (PLAIN)	EXAMPLES OF WHAT IT CATCHES
Content	The words themselves	Known scam scripts (CRA arrest, pig-butchering...), pressure language, "send gift cards"
Artifact	The links & things attached	Lookalike domains, homoglyphs, raw-IP links, fake bank pages, crypto wallets
Authenticity	Voice/video realness	AI-synthesis likelihood (always shown as a probability, never certainty)
Network	What the community reported	Numbers/domains other Canadians already flagged — protection in minutes
Anomaly	What's odd <i>for you</i>	A \$4,200 transfer at midnight to a new payee — abnormal even with no known script

SECTION 11

What happens when you check something



REAL EXAMPLE (ACTUALLY RUNS TODAY)

Paste *"INTERAC: your transfer is pending, accept at <http://interac-secure-deposit.xyz/login>"* → Score **1 / Dangerous**. Reasons: lookalike of Interac · risky ".xyz" domain · flagged by 41 community reports. One tap to block & report. The whole thing happens on the phone in milliseconds.

This is not a mockup — it is the shipped engine. The reasoning "ledger" (which signal contributed how much) is shown in-product, which is what makes the calibrated-honesty moat tangible.

The 30 features, one by one

BUILT works today **DEMO** needs telephony/OS/partner **PLANNED** roadmap · ★ = nobody else offers it · costs are one-time build, CAD.

Pillar 1 — Detect & Decide

Trust Engine & Score ★ **BUILT** P0 \$70–100K

IN PLAIN WORDS
The brain. Turns anything into one 1–100 score with reasons.

HOW IT WORKS
5 signal families → transparent combiner → score + top-3 reasons + ledger + uncertainty, localized to 12 languages, every decision logged for the public calibration report.

Ask KinShield ★ **BUILT** P0 \$45–65K

IN PLAIN WORDS
The "paste anything" box. Text, email, ad, job offer, voicemail, screenshot, QR — instant verdict.

HOW IT WORKS
Share-sheet + keyboard extension + widget → on-device OCR → language ID → engine → action buttons. Proactive: offers a check when another shield sees something odd.

SMS / RCS Shield **DEMO** P0 \$40–60K

IN PLAIN WORDS
Scam texts get filtered to junk before you ever see them.

HOW IT WORKS
iOS `ILMessageFilterExtension` classifies *unknown-sender* messages in a sandbox; Android takes the default-SMS/spam role. On-device multilingual classifier; one-tap report to the network.

Link & QR Checker

BUILT

P0

\$30–45K

IN PLAIN WORDS

Paste a link or scan a QR poster; know in 2 seconds if it's fake.

HOW IT WORKS

Domain age/WHOIS/certificate-transparency + homoglyph/IP/lookalike forensics vs a Canadian brand whitelist; real in-app camera scanner (built); MV3 browser extension.

VerifyCA ★

DEMO

P1

\$45–65K

IN PLAIN WORDS

Check if an "investment firm" or "employer" is actually registered & real in Canada.

HOW IT WORKS

Scraper/API layer over CSA registration + provincial regulators + corporate registries, freshness-monitored, with fuzzy entity resolution.

Multilingual Defense ★

BUILT

P0-P1

\$90–150K

IN PLAIN WORDS

Works in 12 languages where Canadian losses concentrate — not just English.

HOW IT WORKS

Multilingual ASR + per-language scam corpora with settlement-agency partners; a language only ships when paid native speakers validate it (the rest fall back, honestly labeled).

Zero-Day Anomaly Signals

DEMO

P1

\$45–70K

IN PLAIN WORDS

Catches brand-new scams nobody's seen yet, because they're weird *for you*.

HOW IT WORKS

Per-user isolation-forest baselines → anomaly score into the engine; weekly recalibration.

Pillar 2 — Calls & Voice

The KinShield Line ★

DEMO

P1

\$130–190K

IN PLAIN WORDS

An AI receptionist answers unknown callers, figures out if they're a scam, and only patches the real ones through — on any phone.

HOW IT WORKS

Conditional call-forwarding of unknown callers → cloud number (Twilio/Telnyx) → streaming ASR → Trust Engine → policy engine → warm transfer with a whisper announcement. Lives at the network layer, so it's handset-independent.

WHY IT MATTERS

This is how we defend calls without needing forbidden device access — the single most important "how do we make it real" answer.

Caller Intelligence

BUILT

P0

\$35–55K

IN PLAIN WORDS

Flagged numbers get labelled or silenced before your phone rings.

HOW IT WORKS

iOS Call Directory + Live Caller ID Lookup (privacy-preserving); Android caller-ID role; STIR/SHAKEN attestation badge. (Number lookup is live in the app.)

Assist Mode ★

DEMO

P0–P1

\$55–80K

IN PLAIN WORDS

On a suspicious call, put it on speaker and tap Assist — it coaches you live: "This is the CRA-arrest script. Don't give the code."

HOW IT WORKS

On-device streaming transcription via the microphone (permitted) → script classifier → spoken prompts. Lawful one-party consent; nothing leaves the phone.

Voice Lock

DEMO

P1

\$40–60K

IN PLAIN WORDS

Verifies a "family member's" voice against an enrolled voiceprint — catches the cloned-voice trick.

HOW IT WORKS

On-device speaker embeddings (never raw audio) + liveness + anti-replay; Quebec CAI biometric declaration before launch there.

Honeypot Decoy Line ★

PLANNED

P2

\$75–115K

IN PLAIN WORDS

"The app that fights back" — decoy AI grandparents waste scammers' time and steal their playbook.

HOW IT WORKS

Decoy numbers + persona library on our own recording/consent footing; harvested scripts & numbers feed the corpus. Only behind counsel review; never a user's voice. (Precedent: O2 "Daisy", Apate.ai.)

Pillar 3 — Family & Seniors

Family Circle ★

BUILT

P0

\$60–85K

IN PLAIN WORDS

Protect your whole family under one shield — up to 12 people, across iPhone & Android. The senior never manages a password.

HOW IT WORKS

Household graph + roles/permissions consumed by every shield; cross-platform push alerts; family-share key recovery; audited approval log. **The retention & virality engine.**

Senior Mode ★

BUILT

P0

in Family Circle

IN PLAIN WORDS

A big-button, large-text, talk-to-me version that every feature renders into.

HOW IT WORKS

An alternate presentation layer (Help · Family · Check Call · Check Message); spoken verdicts; an acceptance gate on every feature ("if it doesn't work for a 78-year-old in Punjabi, it isn't done").

Scam Drill ★

BUILT

P1

\$30–45K

IN PLAIN WORDS

A friendly game: spot the scam, score points, build the family reflex. Like a fire drill.

HOW IT WORKS

Sanitized templates from the live corpus + reveal-and-teach + household scoreboard; strict ethics rails (never real fear hooks).

Briefing & Heat Map

BUILT

P1

\$25–40K

IN PLAIN WORDS

A weekly "what's going around near you" + a live map of scams across Canada.

HOW IT WORKS

Aggregation over the network plane, city-level only with k-anonymity thresholds; shareable cards that double as marketing.

Pillar 4 — Money & Business

Check Before You Send ★

BUILT

P1

\$30–45K

IN PLAIN WORDS

A 10-second check before money moves: who's the recipient, and three pressure questions.

HOW IT WORKS

Recipient lookup (community + CAFC + crypto-intel) + the three questions (did you initiate? met them? rushed?) + a cooling-off nudge + optional family ping.

Wire Guard (SMB) ★

DEMO

P2

\$70–100K

IN PLAIN WORDS

For businesses: any big payment change needs a live selfie + spoken codeword + voice match — the deepfake can't fake the channel it doesn't control.

HOW IT WORKS

Out-of-band release ceremony + policy engine per org + webhooks (QuickBooks/Stripe) + signed audit log. **Procedure beats detection.**

Meeting Assist ★

PLANNED

P2

\$95–145K

IN PLAIN WORDS

A second opinion on high-stakes video calls — flags deepfake oddities and forces out-of-band verification for new-payee requests.

HOW IT WORKS

Consented screen-capture → frame sampling → detection ensemble → advisory flags (never verdicts) + a hook into Wire Guard.

KinShield SDK ★ (the Bill C-15 product)

DEMO

P1 → P2

\$140–200K

IN PLAIN WORDS

The same brain, sold to banks & credit unions so they can meet the new law.

HOW IT WORKS

REST/GraphQL gateway + tenant isolation + white-label control kit + threat-intel feed + FCAC pipeline + a compliance artifact pack (SOC 2, B-13, residency).

Insurance Benefit

PLANNED

P2

\$80–130K setup

IN PLAIN WORDS

If you're defrauded on a paid tier, a reimbursement benefit — with claims already evidenced by Incident Mode.

HOW IT WORKS

Insurer/MGA partnership; realistic caps (\$10–25K); claims API fed by Incident Mode's hashed evidence. "Designed fast," never "instant."

Pillar 5 — Privacy & Exposure

Zero-Knowledge Vault ★

DEMO

P0

\$70–100K

IN PLAIN WORDS

Your private data is locked on your phone before it's stored — even we can't read it, and a subpoena gets nothing.

HOW IT WORKS

Client-side encryption (libsodium) + Argon2id keys + Shamir family recovery shares + Nitro Enclave for rare server ops + a published external crypto audit.

Exposure Sweep

DEMO

P1

\$60–90K

IN PLAIN WORDS

Finds what scammers already know about you (breaches, data brokers) and shuts it down.

HOW IT WORKS

Breach feeds + broker scrapers + a legal-demand templating engine (statutory Law-25 deletions, not polite opt-outs) + Shadow Contacts; findings feed the engine.

Passkey & Hardening Coach

BUILT

P1

\$25–40K

IN PLAIN WORDS

A gentle checklist that locks down your accounts, with a family "hardening score."

HOW IT WORKS

Curated playbooks per institution/carrier + deep links into settings + ties into Exposure Sweep findings.

Pillar 6 — When It Goes Wrong

Incident Mode ★

BUILT

P0

\$50–75K

IN PLAIN WORDS

The panic button. Compresses the catastrophic first hours into a calm, guided checklist.

HOW IT WORKS

Flow engine + institution playbooks + evidence vault + pre-filled CAFC/police + scripted bank calls + the recovery-scam pre-arm. Automation deepens as bank SDKs land.

Crisis Vault ★

DEMO

P1

\$50–75K

IN PLAIN WORDS

Sextortion help for the four hours that decide everything: de-escalate, preserve evidence, take images down, tell one trusted person.

HOW IT WORKS

Crisis flow + encrypted evidence + StopNCII/TakeItDown hash removal; scripts vetted with mental-health partners; safeguarding gates.

Recovery Line (human) ★

PLANNED

P1

\$30K + \$15–30K/mo

IN PLAIN WORDS

A trained, trauma-informed human, in your language, who joins your bank calls and checks in for 3 months.

HOW IT WORKS

Hiring/training pipeline + QA'd multilingual scripts + case management tied to Incident Mode + strict access controls. An operations moat.

Long-Con Radar ★

BUILT

P1

\$60–85K

IN PLAIN WORDS

For romance & "pig-butcher" investment scams that unfold over weeks — it watches the story progress and steps in before the big ask.

HOW IT WORKS

Per-thread encrypted conversation state + stage classification over message sequences + staged intervention co-designed with victim orgs.

Pillar 7 — Community & Intel

Community Threat Network

BUILT

P0

\$45–70K

IN PLAIN WORDS

One person's report protects everyone else within minutes.

HOW IT WORKS

Consented reports → human-reviewed ingestion → graph propagation → local caches; reputation decay + contributor karma (resists poisoning); aggregate forwarding to the CAFC.

Email Guard

DEMO

P1

\$55–80K

IN PLAIN WORDS

Scans your inbox (with permission) for phishing, and explains why — not just spam-folding.

HOW IT WORKS

Read-only revocable OAuth + SPF/DKIM/DMARC forensics + content model + link checking; bodies discarded after the verdict.

SECTION 13

The data flywheel — how it gets smarter

More users

→

More consented reports

→

Richer Canadian corpus (12 languages)

→

Better detection

→

Fewer losses → trust → stories

→

More users

+ Honeypot harvests scripts

IN PLAIN WORDS

Every scam KinShield sees makes it better at stopping the next one — and the honeypot actively steals scammers' playbooks. The data gets better as we grow, and it's data no competitor can buy because it's Canadian, multilingual, and consent-clean. **This is the compounding engine.**

BOOK IV

The technology

For engineers — each with a plain-words summary so non-technical readers can follow.


SECTION 14

System architecture & the split data-plane


IN PLAIN WORDS

We split data into two buckets that never mix. **Your stuff** (messages, evidence, family) is locked so even we can't read it. **The scammer's stuff** (bad numbers, fake domains) is shared to protect everyone — only when you choose to report it.

PERSONAL PLANE — encrypted, yours, unreadable to us



NETWORK PLANE — consented attacker artifacts, trains the models



The Trust Engine pipeline


HOW IT WORKS

Distilled models run **on-device** (privacy + speed); heavy authenticity ensembles run on a GPU service. The on-device rule corpus (shipped today) is the legible baseline the production models are trained to agree with — which is what lets us publish calibration honestly.

STAGE	TECH	LATENCY BUDGET
Transcription (audio)	Whisper-class ASR, on-device where possible	—
Content classification	Fine-tuned multilingual LLM + rule corpus	<2s text
Artifact forensics	WHOIS/CT/homoglyph/IP/lookalike	<2s links
Authenticity	AASIST/RawNet-class audio ensemble + visual checks	<5s audio · <8s video
Network	Graph lookups + crypto-intel (Chainalysis-class)	—
Anomaly	Per-user isolation forests	—
Combiner	Calibrated noisy-OR (today) → gradient-boosted (prod)	—

SECTION 16

The Line — how cloud call-screening works



IN PLAIN WORDS

Your phone forwards calls from strangers to a KinShield phone number in the cloud. There, a polite AI asks who's calling and listens for scam patterns, then either rings you through (with a heads-up) or blocks it. Because this happens in the cloud, it works on **every** phone — and never needs to touch your call audio on the device.

THE COST REALITY

Per-minute telephony is the biggest running cost. Mitigations: default to unknown-callers-only, per-account minute budgets, dual-provider failover, and gate the Line to paid tiers.

The phone problem — what Apple/Android allow, and our solution

IN PLAIN WORDS

People worry "Apple won't let you read texts or calls." True — but Apple *does* allow the specific, safe things we need, and for the hardest part (call screening) we go around the phone via the cloud. Every capability has a legal, approvable path.

CAPABILITY	IOS PATH	ANDROID PATH
Filter scam texts	SMS Filter Extension — unknown senders only, sandboxed	Default-SMS / spam role (+ Permissions Declaration)
Label/block bad numbers	Call Directory + Live Caller ID Lookup	Caller-ID / spam role; CallScreeningService
Screen the actual call	Network layer — cloud conditional-forwarding (handset-independent)	
Coach you live on a call	Assist Mode via the microphone, on-device (one-party consent)	

WHY IT MATTERS

These limits hit *every* competitor equally — they're not a weakness. Our value was never privileged device access; it's judgment, family, Canadian/multilingual context, and recovery.

Security, privacy & compliance

AREA	APPROACH	WHEN
Encryption	Client-side E2EE of the personal plane; external crypto audit published	P0
Residency	AWS ca-central-1 exclusively — a Canadian sales feature	P0
SOC 2	Type I → Type II (a prerequisite for bank/SDK sales)	P1 → P2
Privacy law	PIPEDA PIAs; Quebec Law 25 / CAI biometric declarations before Voice Lock	P0–P1
Telecom	CRTC / STIR-SHAKEN fluency; honeypot interception review by counsel	P1–P2
Banking	OSFI B-13 mapping for the SDK	P1 → P2

WHY IT MATTERS
Every compliance artifact doubles as B2B sales collateral. Compliance isn't overhead here — it's Moat #3, productized.

BOOK V

Building it — cost, time, people

What it costs, what blocks each feature, the order to build in, and who to hire.

SECTION 19

What every feature costs & how long it takes

One-time build labour, CAD, realistic (senior blended rate), including QA — not optimistic minimums.

FEATURE	COST	BUILD	PHASE
Trust Engine & Score	\$70–100K	8–10 wk	P0
Ask KinShield	\$45–65K	6–8 wk	P0
SMS / RCS Shield	\$40–60K	6–8 wk	P0
Link & QR Checker	\$30–45K	4–6 wk	P0
Caller Intelligence	\$35–55K	6–8 wk	P0
Family Circle (+ Senior Mode)	\$60–85K	10–12 wk	P0
Community Threat Network	\$45–70K	8–10 wk	P0
Incident Mode	\$50–75K	8–10 wk	P0
Zero-Knowledge Vault (+ audit)	\$70–100K	12–14 wk	P0
Multilingual (to 12)	\$90–150K	staged	P0–P1
The KinShield Line	\$130–190K	14–18 wk	P1
Assist Mode	\$55–80K	8–10 wk	P0–P1
Check Before You Send	\$30–45K	5–7 wk	P1
Crisis Vault	\$50–75K	9–12 wk	P1
Email Guard	\$55–80K	8–10 wk	P1
Exposure Sweep	\$60–90K	10–12 wk	P1
Long-Con Radar	\$60–85K	10–12 wk	P1
VerifyCA	\$45–65K	6–8 wk	P1
Voice Lock	\$40–60K	7–9 wk	P1
Zero-Day · Passkey · Scam Drill · Briefing	\$125–195K	—	P1
Recovery Line (setup + opex)	\$30K + \$15–30K/mo	by mo 12–15	P1
KinShield SDK (to pilot)	\$140–200K	16–20 wk	P1 → P2
Honeypot · Insurance · Meeting Assist · Wire Guard	\$320–490K	—	P2
PHASE TOTALS	P0 \$445–650K (BOOTSTRAP CUT \$250–350K) · P1 \$840K–1.23M · P2 \$420–620K · TOTAL ≈ \$1.7–2.5M OVER ~3 YRS		

SECTION 20

The obstacle in each feature & how it's resolved

FEATURE	OBSTACLE	RESOLUTION
SMS Shield	iOS only sees unknown senders, no free network; Android needs a Play declaration (rejection risk)	On-device model (no network); private-relay fallback; demo video + justification — peers pass this routinely
The Line	Per-minute cost is the top opex lever; per-carrier forwarding; transfer latency	Unknown-callers-only default; minute budgets; dual-CPaaS; per-carrier wizard; streaming ASR <2s
Caller Intelligence	Apple gates the Live Caller ID entitlement; blocklists go stale fast	File entitlement week one; bloom-filter cache; lean on behaviour (the Line) not just number
Assist Mode	On-device ASR battery/thermal; noisy speakerphone	Distilled streaming ASR; thermal caps; degrade to keyword spotting
Voice Lock	Biometric law (CAI); spoof/replay; senior enrollment friction	CAI declaration; store embeddings not audio; liveness + anti-replay; senior-designed UX
Honeypot	Interception/recording & CRTC law; never a user's voice	Counsel gate; own DIDs & consent footing; strict multi-report thresholds (precedent exists)
Family Circle	Monitoring a parent → consent & coercive-control risk	Consent ceremonies; transparency to the protected member; no covert mode; audit log
Incident Mode	No APIs to freeze accounts or file reports automatically	Pre-filled deep links + sequenced playbooks (content ops); automation deepens with bank SDKs
Crisis Vault	Safeguarding, minors, mental-health liability	Mental-health-vetted scripts; age gating; StopNCII/TakeItDown; mandatory safeguarding review
Recovery Line	Hiring trauma-informed multilingual staff; cost; advice liability	QA'd scripts; strict scope (navigation, not legal advice); paid-tier gated; scales with revenue
Insurance	12+ month underwriter negotiation; actuarial risk; OSFI	MGA partner; realistic caps; pre-evidenced claims; honest "designed fast"
VerifyCA	13 provincial registries, no APIs, entity resolution	Scrapers + freshness monitoring + fuzzy matching + OCR
SDK	6–12 month bank sales; SOC 2 prerequisite	Credit unions first; compliance pack; Big-5 in Yr 3 on references
Trust Engine / ML	Scarce Canadian multilingual data; drift; adversarial evasion	Corpus + community + honeypot flywheel; drift monitoring; published calibration
ZK Vault	Encryption vs senior recovery; audit cost	Shamir family shares; external audit (~\$20–30K) published before marketing the claim

SECTION 21

The four-stage roadmap, step by step

Stage 0 · Validate — 0–3 mo · \$8–18K

Do: clear the name; incorporate + founder IP assignment; 100 discovery interviews; per-language waitlist + small ads; a manual concierge MVP for 50 families (builds the first labeled corpus by hand); 1 partnership LOI; lawyer sanity check; apply IRAP/SR&ED.

Challenge: a deck feels like progress. **Fix:** force the concierge test; if <15% call it serious or no repeat sends, pivot for \$5K — don't build.

Exit gate: 1,500 waitlist · 1 LOI · repeat concierge use · name cleared.

Stage 1 · MVP — 3–9 mo · \$90–350K

Build (zero app-store risk): SMS Shield, Link/QR, Ask, Trust Score (4 langs), Family Circle + Senior Mode, Incident Mode, Community reporting.

Challenge: native hooks + first ML model are hard. **Fix:** hire 2 specialist contractors only for those; file Apple entitlement & Play declaration in month 3.

Exit gate: 10K installs · 500+ paying · D30 ≥25% · 1 documented save · CAC <\$20.

Stage 2 · Seed & scale — 9–24 mo · \$1.5–2.5M

Ship in order: call layers + Exposure Sweep → the Line + VerifyCA + Check-Before-Send → Voice Lock + Crisis Vault + **SDK pilot #1 (credit union)** → Wire Guard + languages 9–12 + SDK pilot #2 + insurance talks begin.

Challenge: telephony cost; SOC 2 gate; insurance lead time. **Fix:** dual-CPaaS + paid gating; start SOC 2 Type I at month 9; begin underwriting in P1.

Exit gate: \$1M+ ARR · 2 SDK pilots · churn <3.5%/mo · SOC 2 Type I.

Stage 3 · Series A & category — yr 2–4 · \$8–15M

Do: bank SDK revenue; insurance live; Honeypot + Meeting Assist; multi-tenant SDK; Big-5 / credit-union-central deals on SOC 2 Type II + references.

Exit gate: \$5–10M ARR · a Big-5 or CU-central deal · path to profitability.

Who to hire, and when

STAGE	TEAM	NOTES
0 Validate	Founder(s) + community validators (paid)	Evenings/weekends; no full-time hires
1 MVP	Founder full-time + 1 senior mobile contractor + 1 ML contractor + part-time designer	Hybrid: build what you can, hire only the 2 hard narrow pieces
2 Seed	+ senior backend/telephony eng, ML eng, mobile eng, head of community partnerships, part-time compliance lead, 2 support/recovery staff	7–9 people incl. founder; hire in this order
3 Series A	+ enterprise sales, customer success, more ML/recovery, finance/ops, security lead	Scale the recovery org & the SDK go-to-market

IN PLAIN WORDS

You don't hire a big team early. You stay tiny and scrappy through the MVP, using contractors for the two genuinely hard pieces (native phone hooks + the first AI model), and only build a real team once paying customers prove it works.

BOOK VI

The business

Pricing, how we get customers cheaply, the five-year model, and the money behind it.

Pricing & packaging

TIER	PRICE	WHO BUYS IT & WHY
Free	\$0	Everyone — distribution + network data. SMS Shield, basic checks, self-serve Incident Mode.
Family	\$9.99/mo	The wedge. The adult child buys peace of mind for a parent — emotional, cheap. The Line, Family Circle, Crisis Vault, Long-Con, Vault.
Premium	\$19.99/mo	SMBs & high-exposure users — adds the human Recovery Line, Wire Guard, Voice Lock, insurance.
Enterprise SDK	\$50–250K/yr	Banks & credit unions meeting Bill C-15. The institutional revenue engine.

IN PLAIN WORDS

The free version spreads the app and feeds the data. The \$9.99 family plan is priced like a coffee-a-week so an adult child says yes instantly for a parent. Premium and the bank SDK are where the real margin is.

Go-to-market — how we get users cheaply

CHANNEL	HOW IT WORKS	WHY IT'S CHEAP
Community partnerships	Settlement agencies, seniors' orgs, libraries, credit unions co-launch & co-design	Borrowed trust; CAC <\$20; also the source of the multilingual corpus
Family virality	One buyer enrolls up to 12 — parents recruit kids, kids recruit siblings	Near-zero CAC; the family graph is the loop
Earned PR	Real, consented "money saved" stories; published calibration journalists can trust	Free press; the honesty moat earns coverage
Multilingual content / SEO	Scam alerts, shareable briefing cards, "is this a scam?" guides per language	Organic; nobody else publishes in these languages
B2B2C distribution	Carriers, banks, employers, insurers push it to their users	Their channel, our product
Paid ads	Targeted by segment & language — for measurement, not the main channel	Used to find product-market fit signals, kept small

The five-year financial model

Conservative planning estimate (CAD), post-MVP launch. Not a guarantee — a base case to pressure-test.

METRIC	YR 1	YR 2	YR 3	YR 4	YR 5
Installs (cumulative)	25K	120K	400K	900K	1.8M
Paying households	2.0K	12K	45K	110K	220K
Free → paid conversion	8%	10%	11%	12%	12%
Consumer ARR	\$0.28M	\$1.68M	\$6.3M	\$15.4M	\$30.8M
SDK customers	1	3	8	15	25
SDK ARR	\$0.08M	\$0.3M	\$1.2M	\$3.0M	\$6.0M
Insurance / telecom ARR	—	—	\$0.3M	\$1.0M	\$3.0M
TOTAL ARR	\$0.36M	\$2.0M	\$7.8M	\$19.4M	\$39.8M
Gross margin	60%	70%	78%	82%	84%
Headcount	4	9	22	45	80
Capital stage	pre-seed \$0.3–0.5M	seed \$1.5–2.5M	Series A \$8–15M	(A)	B optional

WHY IT MATTERS

Even capturing <1.5% of Canadian households plus a modest SDK book reaches ~\$40M ARR by Year 5 — and that's **before** the US and global diaspora markets (Book VII), which are far larger.

Unit economics — does each customer make money?

METRIC	VALUE	NOTE
ARPU (paying household)	~\$140 / yr	Blended Family + Premium
Gross margin at scale	78–84%	Line gated to paid; minute controls
CAC (blended)	\$20–40	Partnership + family-viral led
Annual churn	~30% → improving	Family graph lowers it over time
LTV (gross profit)	~\$350–400	~3-yr lifetime × ~\$112 gross profit/yr
LTV : CAC	~10:1	Healthy; payback <3 months

IN PLAIN WORDS

Each paying family costs ~\$20–40 to acquire and is worth ~\$350+ over time — roughly a 10× return — and pays us back within three months. That's the engine that lets marketing spend grow safely.

Funding & free government money

STAGE	RAISE	INSTRUMENT / SOURCE
0 Validate	\$8–18K	Founder savings
1 MVP	\$0.3–0.5M	Friends/angels SAFE (bootstrap path possible at \$90–160K)
2 Seed	\$1.5–2.5M	SAFE or priced seed; defensible cap \$8–14M on traction + the C-15 narrative
3 Series A	\$8–15M	Institutional VC on \$5–10M ARR + a bank/CU-central deal

IN PLAIN WORDS — FREE MONEY

SR&ED refunds roughly **35% of eligible R&D salaries/contractors** for a Canadian-controlled company (often \$25–60K back on the MVP budget), and **NRC IRAP** co-funds technical work. File from the first line of R&D code — it's the difference between 6 and 10 months of runway.

BOOK VII

The future — how this grows huge

The recommendations you asked for: where KinShield goes after the first 30 features.

Growth levers (next 12 months)

- **Widen languages fast** — each validated language opens a whole diaspora community with near-zero competition.
- **Sign the first carrier pilot** — a Rogers/Bell/Telus network-level deal puts the Line in front of millions.
- **Land the first credit-union SDK** — the reference customer that unlocks the rest of the C-15 buyers.
- **Launch a family-referral program** — formalize the viral loop ("protect your parents, get a month free").
- **Build the PR engine** — a steady cadence of consented "real save" stories + the quarterly calibration report.

Future features (beyond the 30)

IDEA	WHAT IT ADDS
KinShield Assistant (AI agent)	A conversational guardian that proactively watches your channels and warns you before you ask — the natural evolution of Ask KinShield into an always-on agent
Bank-account connection (consented)	Real-time transaction risk + the consumer side of Bill C-15 controls (cool-off, limits) — deepens the money-defence and the SDK story
Teen & kids mode	Sextortion, gaming & social scams aimed at minors — a whole new buyer (parents of teens) on top of seniors
Estate / Power-of-Attorney mode	Safely co-manage a declining parent's finances — a sensitive, high-trust, sticky use case
Remittance / send-abroad guard	Protect newcomers sending money home from remittance & emergency scams
Real-time deepfake call detection	As on-device models mature, live synthesis flags during a call — the long-term Pindrop-at-consumer-scale play
Full desktop & browser guard	Marketplace/classifieds/dating "check before you trust" everywhere you browse
Voice-assistant integration	"Hey Siri, is this a scam?" — distribution through the assistants people already use

New markets & new countries

IN PLAIN WORDS

The hard part — speaking many languages and understanding diaspora scams — is exactly what makes us travel well. The same engine works anywhere; only the local registries, brands and laws change.

MARKET	WHY IT FITS	WHAT CHANGES
USA (the big one)	10× the market; Aura's turf but no multilingual/family depth; huge Spanish-first & Asian diaspora	US registries, brands, FTC/CFPB framing; Spanish-first
UK · Australia (Commonwealth)	Similar regulation, scam patterns & English+diaspora mix; Australia's scam codes mirror C-15	Local regulators, carrier deals
EU	Strong privacy law favours our E2EE posture; many languages	GDPR, per-country registries

WHY IT MATTERS

"Canada's fraud-defence app" is the wedge. "The fraud-defence layer for every multilingual family on earth" is the company. Win Canada first to prove the model, then export the playbook.

New revenue: telecom, insurance, government, B2B

STREAM	THE PLAY
Telecom partnerships	Carriers embed the Line network-side (Hiya's model) — but the consumer relationship stays ours. Revenue share + reach.
Insurance-as-a-product	Become an MGA: detection telemetry lowers actuarial risk → cheaper premiums → a flywheel competitors can't price against.
Government & CAFC	Aggregate, consented threat data & tooling for national anti-fraud efforts — aligned with the National Anti-Fraud Strategy.
B2B employee protection	Wire Guard + Meeting Assist sold to companies to stop BEC/deepfake-CEO fraud — a direct enterprise SaaS line.
Threat-intel feed	Sell the consent-clean Canadian artifact feed to other security vendors & banks.

SECTION 32

Becoming a platform — the trust layer

IN PLAIN WORDS

Open the Trust Engine to other apps. A marketplace, dating app, classifieds site or crypto exchange can embed "Check before you trust" with one API call — KinShield becomes the **"Stripe for trust"**: the layer everyone calls to ask "is this safe?"

WHY IT MATTERS

Every integration feeds the network plane and widens the moat. The consumer app builds the brand and the data; the platform API turns that data into a category-defining utility with network effects.

Who buys us — the exit landscape

ACQUIRER TYPE	EXAMPLES	WHY THEY'D WANT IT
Consumer-security	Gen Digital (Norton/Avast/LifeLock), Aura, McAfee	Instant Canadian + multilingual + family capability
Banks	RBC, TD, a credit-union central	Build vs buy their Bill C-15 obligation; the SDK + recovery
Telecoms	Rogers, Bell, Telus	Network-level scam defence + a consumer relationship
Payments / networks	Mastercard, Visa, Interac	Transaction-fraud + the trust API
Platforms	Google, Apple, Microsoft	Acqui-hire / feature; the multilingual corpus

IN PLAIN WORDS

Many natural buyers, each of whom would pay a premium because KinShield gives them something they can't quickly build: the data, the languages, the family graph, and the recovery operation. Or — at enough scale — it's an IPO-class category leader. Comparable players (Aura, Hiya, Pindrop) have raised and exited at strong multiples.

BOOK VIII

Risks & realities

An honest accounting — the risks, the things founders forget, the hard investor questions, and the ethics.

The risk register

RISK	LIKELIHOOD	IMPACT	MITIGATION
App Store / Play rejection of extensions	Med	High	Sanctioned APIs only; file early; demo build + justification; peers ship these
Telephony cost erodes margins	Med	Med	Unknown-callers-only; minute budgets; gate to paid; dual-CPaaS
Scammers adapt to detection	High	Med	Anomaly net + honeypot/community flywheel; published calibration so it's visible
Slow consumer willingness to pay	Med	High	Adult-child-buys-for-parent emotional wedge; cheap family tier; free distribution
Privacy regulator action	Low	High	Compliance is the product (Moat 3); PIAs, residency, consent-clean network
Brand-damaging overclaim	Low	High	Claims discipline; never publish a number without benchmark/date/failures
Key-person / talent scarcity	Med	Med	Hybrid build; boring/hireable stack; IP assignment from commit one
Giants bundle a "good enough" version	Med	Med	Orchestration + family + Canada + recovery is a layer they won't run; partner where useful
Insurance underwriting falls through	Med	Low	Not core to the wedge; multiple MGA partners; ships in P2 only
Capital market tightens	Med	Med	Phase gates; SR&ED/IRAP non-dilutive; bootstrap path keeps MVP alive

What founders forget

OFTEN FORGOTTEN	WHAT TO DO
Native extension review lead time	SMS Filter, Call Directory & Live Caller ID are separate iOS targets with their own review — file in week one
Content operations	Bank scripts, registry lists & playbooks rot — budget a content-ops function, not just engineering
Abuse & cost controls on the Line	One abused account can blow the telephony budget — caps & monitoring from day one
Accessibility as a gate	Senior Mode is an acceptance criterion on every feature, validated in paid usability studies
Data poisoning & model drift	Contributor karma + human review at ingestion; drift monitoring + per-language calibration
Coercive-control consent edge cases	Explicit consent ceremonies; transparency to the protected member; no covert monitoring
Recovery as a real org	A hiring/training/QA pipeline, not a ticket queue — and the source of the brand's stories
Per-cohort unit economics	Track infra- and Line-minutes-per-paying-household, not just blended margin
Entity, trademark & cap-table hygiene	CIPO/NUANS clearance, incorporation, founder IP assignment, clean SAFEs before any marketing
Non-dilutive capital	SR&ED + IRAP from the first line of R&D code
Don't oversell deepfake detection	Detectors generalize poorly — rely on multi-signal + procedure (Wire Guard), and say so
Safeguarding & minors	Sextortion flows touch minors — partner review & safeguarding sign-off are mandatory gates

Investor FAQ — the hard questions

"Won't Apple or Google just do this themselves?"

They commoditize single-channel *detection* — and that helps us, because we build on top of it. They will not run a multilingual family graph, a human recovery line in Punjabi, Canadian institutional context, or a Bill C-15 bank SDK. We sell the layer above them; where they expose APIs, we integrate.

"Telephony costs will kill your margins."

Only if unmanaged. Unknown-callers-only screening, per-account minute budgets, dual-CPaaS, and gating the Line to paid tiers deliver 78–84% gross margins at scale. The Line is a Premium-tier value driver, not a free-tier cost sink.

"Can a small team out-execute giants?"

We're not racing them head-on. We're running the race nobody is running — the orchestration layer — in a market (Canadian, multilingual, family, recovery) they're structurally unsuited to serve.

"Is Canada big enough?"

Consumer alone is a ~\$40M-ARR path by Year 5. Add the C-15 SDK, insurance, telecom — then the US and global diaspora markets (10×+). Canada is the wedge, not the ceiling.

"What if scammers adapt?"

The anomaly net catches zero-days; the community + honeypot flywheel keeps the corpus ahead; and we publish calibration so adaptation is measured, not hidden. Adaptation is the business, not a threat to it.

"Why will people pay for security they hope never to need?"

They don't pay for themselves — the adult child pays for a *parent's* safety. That's an emotional purchase at a coffee-a-week price, which is why the Family tier converts.

Claims discipline & ethics — the moat that costs nothing

- Never publish an accuracy number without the benchmark, the date and the failure cases beside it.
- Never imply protection on a channel the OS doesn't permit.
- Never attribute a breach or failure to a competitor without primary-source verification.
- Never promise insurance outcomes — "designed to be fast," never "instant" or "no proof of loss."
- Never use invented pilot data; never collect what the personal plane can't protect; never ship a language unvalidated by paid native speakers.

THE GENIUS IT RUNS ON

In a market built on "96%" and "99.2%," the company that publishes its calibration curves — misses included — becomes the one journalists cite, regulators trust, banks shortlist, and families recommend. Not more claims than the giants — claims that are **true**.

Reference

Glossary — every term in plain words

Smishing	A scam text message (SMS + phishing).
Pig-butcher	A long romance/investment scam that "fattens" the victim over weeks before the big steal.
Deepfake	An AI-generated fake voice or video of a real person.
Homoglyph	A look-alike character used to disguise a fake web address (e.g. a Cyrillic "a").
STIR/SHAKEN	The phone-network standard that verifies a caller ID isn't spoofed.
E2EE	End-to-end encryption — data only readable on your device, not by the company.
CPaaS	Communications Platform as a Service (Twilio/Telnyx) — cloud phone infrastructure.
MGA	Managing General Agent — a partner that lets you offer insurance without being an insurer.
SOC 2	A security audit standard banks require before buying software.
PIPEDA / Law 25	Canada's federal / Quebec privacy laws (Law 25 has real penalties & deletion rights).
CAFC	Canadian Anti-Fraud Centre — the national fraud-reporting body.
OSFI B-13	The banking regulator's technology/cyber-risk guideline.
SR&ED / IRAP	Canadian programs that refund/fund R&D — non-dilutive money.
BEC	Business Email Compromise — the deepfake-CEO / fake-invoice wire-fraud class.
CAC / LTV / ARPU / ARR	Cost to acquire a customer / lifetime value / avg revenue per user / annual recurring revenue.

The scam scripts the engine already knows

Named "script DNA" the on-device engine detects today, each in multiple languages, plus a brand-agnostic generic-phishing layer:

CRA / SIN-suspension arrest	Interac e-transfer phish	Canada Post / parcel-fee smish
IRCC / immigration-status threat	Grandparent / family emergency	Pig-butcherer (romance investment)
Bank fraud-department impersonation	Tech-support / remote access	Prize / lottery advance-fee
Job-offer / reshipping	Gift-card extraction	Sextortion
Utility disconnection (hydro/gas)	Unpaid toll / 407 ETR	Rental deposit (sight-unseen)
Marketplace overpayment	"Is this you?" account-takeover / OTP	Disaster / charity fraud

+ generic phishing grammar (account locked, unusual sign-in, verify your account, update your card...) and full URL forensics (lookalike, homoglyph, raw-IP, credentials-in-URL, risky TLD, shortener, brand-off-domain).

Sample SDK API (the Bill C-15 product)

```
POST /v1/check
{
  "text": "URGENT: move your balance to a safe account, read me the code",
  "channel": "call_transcript",
  "recipient": "6045550147",
  "language": "en",
  "context": { "userInitiated": false, "amountCAD": 4200 }
}

→ 200 OK
{
  "trustScore": 7,
  "verdict": "dangerous",
  "action": "block_and_report",
  "detectedScript": { "id": "bank_security",
                     "label": "Bank fraud-department impersonation" },
  "ledger": [ { "family": "content", "risk": 0.92 },
              { "family": "network", "risk": 0.86 },
              { "family": "anomaly", "risk": 0.45 } ],
  "engineVersion": "trust-engine/1.0+corpus-2026.06"
}
```

Plus a white-label control kit (cool-off + capability toggles + limits) and the FCAC reporting pipeline.

What's built today & how to run it

PATH	WHAT
mobile/	Native iPhone app (Expo SDK 54) — runs in Expo Go; the flagship
mobile/lib/trust-engine	On-device Trust Engine (shared, pure TypeScript)
src/ (root)	Next.js web app / PWA on the same engine
src/lib/trust-engine	Engine source + 25 unit tests + live calibration harness
docs/	This document + the v1 master doc

iPhone: `cd mobile && npm install && npx expo start` → scan the QR (Expo Go). · **Standalone:** `eas build --profile preview --platform ios`. · **Web:** `npm install && npm run dev`.

The 90-day action checklist

WEEKS	DO
1–2	Incorporate; founder IP assignment; CIPO/NUANS name clearance; domains/socials; open Apple Developer + Play + cloud + CPaaS accounts (the long queues); file the Apple Live Caller ID entitlement.
2–6	100 discovery interviews (seniors, adult children, newcomers, victims, bank/CU risk staff); per-language landing pages + small ads; start the manual concierge MVP for 50 families.
4–8	One settlement-agency/seniors-org LOI; 2–3h with a privacy/fintech lawyer; apply IRAP + set up SR&ED tracking.
8–12	Decide bootstrap vs pre-seed; if the gate passes (1,500 waitlist · 1 LOI · repeat concierge use), begin the MVP build with 2 specialist contractors.